

Theis on:

HOW TO MAKE AI-MODELS DO STUFF!

Leveraging Rust and
Candle to Execute LLM
tasks on IoT



OUR PLAN

- SECTION 1/2 – AI AND RUST
 - Why use Rust for AI? (LLMs)
 - ...Is Rust just better? Not always...
 - What are LLMs?...
 - How Candle supports LLMs
 - Demo 1/2
- SECTION 2/2 – AI DOING THINGS USING RUST
 - Building the AI ecosystem in Rust
 - Task execution and making AI do things
 - The brain: Agent
 - The state: Handling events
 - The state: Routing events
 - The state: Actions
 - Workflow
 - How can we do... Whatever?
 - Demo 2/2
- Further recommendations

SECTION 1/2

AI and Rust

WHY USE RUST FOR AI? (LLM)

| Category | Rust | Python |
|-----------------|---|---|
| Execution Speed | Compiled, faster inference, ideal for real-time AI tasks | Interpreted, slower inference due to GIL and overhead |
| Memory Safety | Precise memory mgmt., preventing leaks mem leak in large models. Safer services | GC risks performance drops and leaks |
| Concurrency | True parallelism, ideal for high-demand AI services, web frameworks and more | GIL limits concurrency, slowing real-time tasks |
| Memory Usage | Low overhead, efficient use of memory for large models. Great for IoT and mobile devices! | Higher mem usage due to dynamic typing and GC |
| Model Loading | Optimized for efficient loading of large AI models | Slower model loading due to abstractions |
| Latency | Low latency, ideal for real-time AI applications | Higher latency due to GIL and overhead |

...IS RUST JUST BETTER? NOT ALWAYS...

... But where it is, it
really shines

| Feature | Rust (Candle) | Python (Transformers) |
|-------------|--|--|
| Performance | ⚡ End-to-end performance, low latency, memory efficiency | ⚡ Optimized with PyTorch/TF, heavy computation tasks (C++ binding) |
| Ecosystem | 🌱 Evolving ecosystem, smaller set of available models for easy setup | 🌐 Large ecosystem, pre-built/trained models, integration with PyTorch/TF |
| Concurrency | ✅ True parallelism, optimized for real-time, multi-threaded tasks | ❌ Limited by Python's GIL, slower real-time inference |
| Use Cases | 🕒 Excellent for real-time inference, resource-constrained envs (edge AI), serverless compute systems | 🚀 Ideal for rapid-prototyping, development and finetuning pre-trained models |
| Ease Of Use | 🔧 Bit more complex setup, requires lower-level mgmt, fewer abstractions (more fun!) | 🎯 Simple API, vast community support, and well documented for most cases |

| Factor | Rust | Python |
|-----------------|-----------|-----------|
| Performance | ■ ■ ■ ■ ■ | ■ ■ ■ ■ ○ |
| Concurrency | ■ ■ ■ ■ ■ | ■ ■ ○ ○ ○ |
| Ecosystem | ■ ■ ■ ○ ○ | ■ ■ ■ ■ ■ |
| Ease Of Use | ■ ■ ■ ○ ○ | ■ ■ ■ ■ ■ |
| Real-Time Tasks | ■ ■ ■ ■ ■ | ■ ■ ○ ○ ○ |

WHAT ARE LLMS?...

...And why do they matter?

- LLMs (like the famous GPT-series and BERT) are built using transformer architectures that process and understand context in vast amounts of text data (See video)
- LLMs are today trained on billions of parameters, allowing them to perform tasks, such as:
 - text generation
 - summarization
 - machine translation
 - question-answering
 - code generation
- Why is this relevant?
 - They are great at context handling, and has become crucial in apps like chatbots, but also more autonomous AI-driven applications, and automated code generation
- They are continuing to grow rapidly, both in terms of performance, weights and knowledge

HOW CANDLE SUPPORTS LLMS

... And how to use it

- Candle efficiently handles large arrays of data, or tensors, which are fundamental in deep learning models, including LLMs
- Built-in support for model inference, Candle simplifies running LLMs like GPT in Rust, eliminating the need for Python
- Candle leverages Rust's performance and concurrency to make AI more accessible for real-time apps on edge devices and servers

... Efficient Tensor Management

| | Using PyTorch | Using Candle |
|------------|--|--|
| Creation | <code>torch.Tensor([[1, 2], [3, 4]])</code> | <code>Tensor::new(&[[1f32, 2.], [3., 4.]], &Device::Cpu)?</code> |
| Creation | <code>torch.zeros((2, 2))</code> | <code>Tensor::zeros((2, 2), DType::F32, &Device::Cpu)?</code> |
| Indexing | <code>tensor[:, :4]</code> | <code>tensor.i(.., ..4)?</code> |
| Operations | <code>tensor.view((2, 2))</code> | <code>tensor.reshape((2, 2))?</code> |
| Operations | <code>a.matmul(b)</code> | <code>a.matmul(&b)?</code> |
| Arithmetic | <code>a + b</code> | <code>&a + &b</code> |
| Device | <code>tensor.to(device="cuda")</code> | <code>tensor.to_device(&Device::new_cuda(0))?</code> |
| Dtype | <code>tensor.to(dtype=torch.float16)</code> | <code>tensor.to_dtype(&DType::F16)?</code> |
| Saving | <code>torch.save({"A": A}, "model.bin")</code> | <code>candle::safetensors::save(&HashMap::from([("A", A)]), "model.safetensors")?</code> |
| Loading | <code>weights = torch.load("model.bin")</code> | <code>candle::safetensors::load("model.safetensors", &device)</code> |



... What do Candle say?

Why should I use Candle?

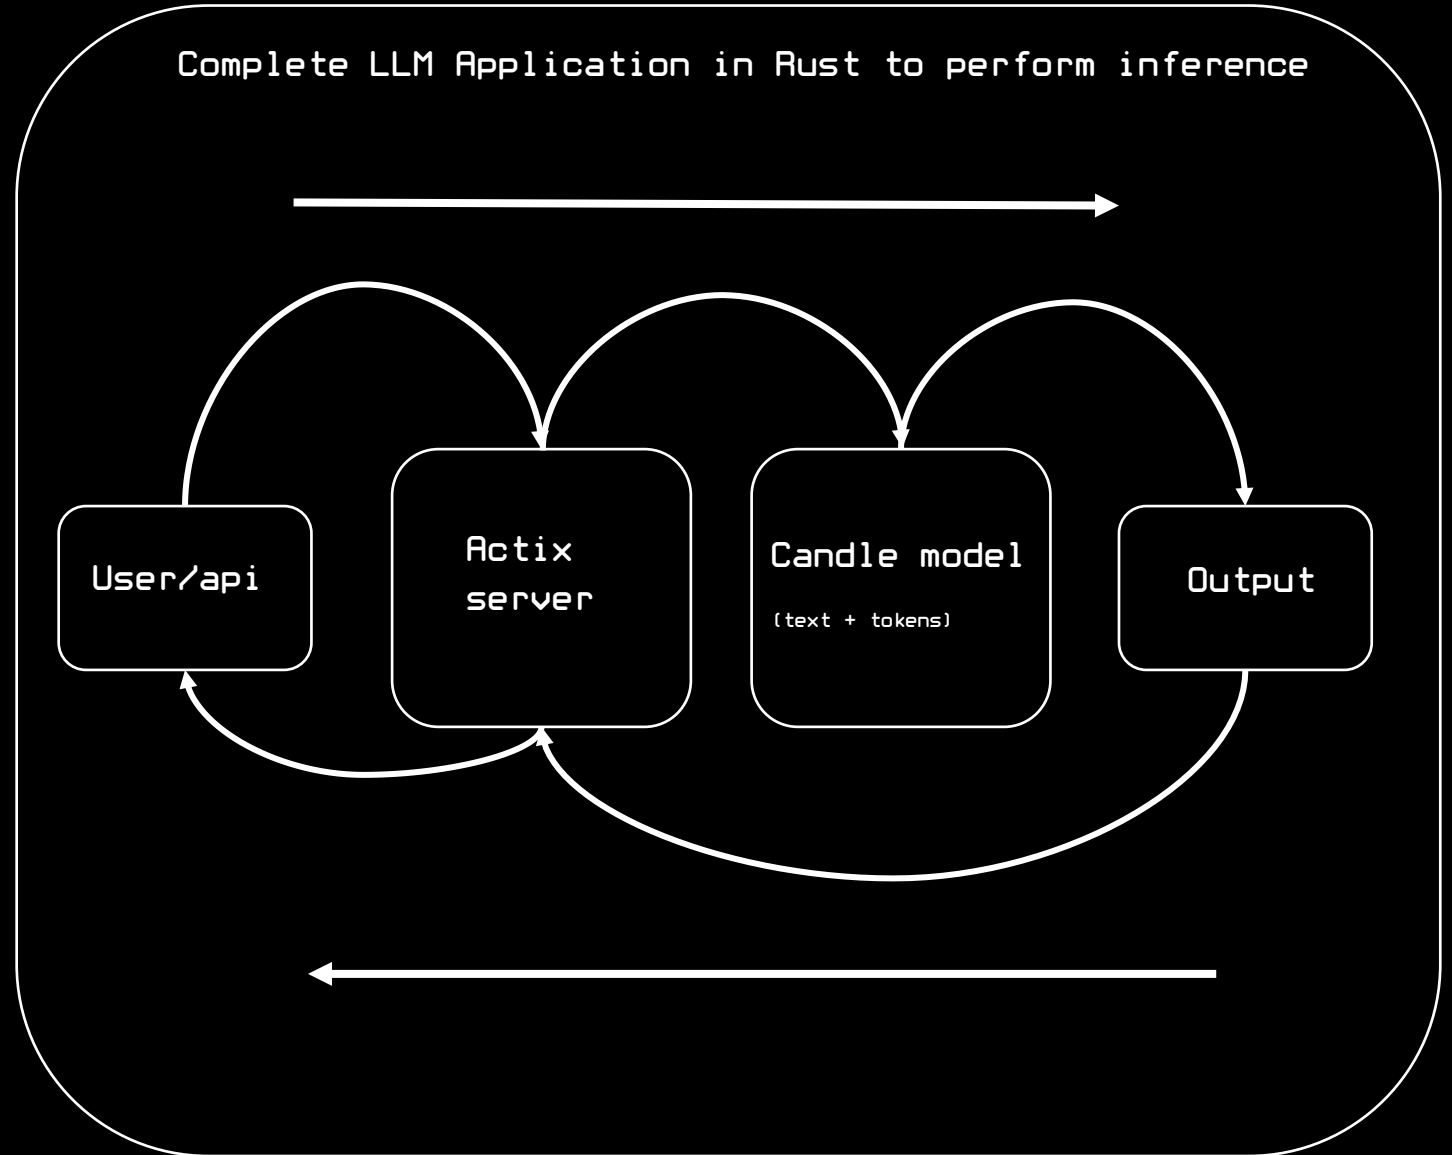
Candle's core goal is to *make serverless inference possible*. Full machine learning frameworks like PyTorch are very large, which makes creating instances on a cluster slow. Candle allows deployment of lightweight binaries.

Secondly, Candle lets you *remove Python* from production workloads. Python overhead can seriously hurt performance, and the [GIL](#) is a notorious source of headaches.

Finally, Rust is cool! A lot of the HF ecosystem already has Rust crates, like [safetensors](#) and [tokenizers](#).

FINALLY, ACTION!

... Lets see how we can
build something interesting



CONTACT

Theis on Discord for
this content

DEMO

1/2

...

SECTION 2/2

Ai doing things
using Rust

CONTACT

Theis on Discord for
this content

HOW CAN WE
DO...

...WHATEVER?

... Lets look at some
cool things we could do

DEMO

2/2

...

FURTHER RECOMMENDATIONS

- Video of tensors:
(<https://www.youtube.com/watch?v=f5liqUk0ZTw>)
- Video of Transformers:
(<https://www.youtube.com/watch?v=wjZofJX0v4M&t=542s>)
- AutoGPT project:
(<https://github.com/Significant-Gravitas/AutoGPT>)
- Candle project:
(<https://github.com/huggingface/candle>)